Girivardhana Reddiar

Data Science, MS Candidate | LinkedIn | reddiarg@gmail.com | +1 (781)-770-1184

Summary

Results-driven data professional with **3 years of experience** with a proven track record of designing and implementing **end-to-end scalable data pipelines** across diverse industries. Adept at translating **complex business needs** into high-performance data architectures, leveraging **Agile methodologies** to drive seamless execution. Recognized for **exceptional client relationship management**, consistently exceeding expectations by delivering **innovative, cost-efficient, and scalable solutions** that enhance **data-driven decision-making**.

Skills

- **Programming & Development**: Python, R, C/C++, SQL, Git, Docker
- **Data tools & Libraries**: Databases (MySQL, Postgres, DynamoDB), Data Warehouse (Redshift, Snowflake), PySpark, FastAPI, Pandas, NumPy, Scikit-learn, Matplotlib, Jupyter Lab & Notebook, AWS Cloud Platform
- **Certifications**: AWS Solutions Architect (Associate), AWS Developer (Associate), Snowflake SnowPro Core
- **Additional**: Machine Learning (regression, classification, clustering), Statistics (inferential and descriptive, hypothesis testing, A/B testing), Microsoft Office 365, JIRA

Experience

Senior Data Engineer, *Quantiphi Analytics*   2023-2024

- Led the end-to-end development of a SaaS-based data platform for an entertainment & sports ventures firm, consolidating **terabytes of consumer data** into an AWS S3 data lake. Enabled custom consumer segmentation and comparative analysis, reducing report generation time from **days to minutes** and **cutting cloud costs by 70%** through optimized data partitioning and AWS EMR Serverless.
- Developed a **multi-layered enterprise data warehouse** on AWS Redshift, integrating disparate tenant databases into a single source of truth for a large HR & talent acquisition firm. Built **6+ dashboards** in QuickSight, enabling **trend analysis** in hiring patterns and future capabilities like fraud detection and engagement analysis.
- Optimized a real-time streaming data pipeline for a cybersecurity firm, identifying bottlenecks and implementing **windowed stream processing** in Kafka, increasing ingestion rates from **3,000 to 250,000 events/second** while maintaining cloud service costs. Designed a **future-state architecture** leveraging AWS Redshift to handle **terabyte-scale data** and improve system resilience.
- Mentored **6 interns**, conducted **20+ technical interviews**, and led **org-wide knowledge-sharing** sessions on reusable components to **enhance team productivity**.

Data Engineer, *Quantiphi Analytics*   2021-2023

- Migrated **terabytes of historical data** from Snowflake to AWS S3, **optimizing storage costs** for an NFT-based analytics startup. Designed an AWS-based data warehouse solution leveraging Redshift for analytics and Timestream for time-series data, **improving scalability** and **cutting cross-cloud data transfer costs**.
- Built an **event-driven preprocessing framework** for a cybersecurity analytics project, parsing network log data in **near real-time** using AWS Lambda, Kinesis Firehose, and EventBridge. Reduced data ingestion delays from **1 day to 2 hours**, improving threat detection efficiency and reducing manual investigation efforts by **85%**

Education

MS, Data Science *University of Massachusetts*   Dartmouth, MA 2024-2026

- Relevant Coursework: Advanced Machine Learning, Advanced Mathematical Statistics, Business Analytics & Data Mining

B Tech, Computer Engineering *NMIMS University*   Mumbai, India 2017-2021